

Integrated Inventory Donation Tracking System for Non-Profits

Design, Validation and Verification Document

Group 14

Insiyah Ujjainwala
Taaliah Ayub
Haley Johnson
Andrew Goetz
Nika Khajehpour
Victor Moucattash
Alyssa Wang

Table of Contents

- 2. Version History..... 3
- 3. Purpose Statement..... 3
- 4. Component Diagram..... 3
- 5. Relationship Between Components and Requirements..... 3
 - 5.1. Backend..... 3
 - 5.2. Frontend..... 4
- 6. Component Descriptions..... 5
 - 6.1. Backend Components..... 5
 - 6.2. Frontend Components..... 5
 - 6.3 Machine Learning Component..... 6
- 7. User Interface..... 7
 - 7.1 Layout, navigation, and visual style..... 7
 - 7.2 Authentication and account actions..... 7
 - 7.3 Home page (overview + forecasting)..... 7
 - 7.4 Inventory page (search + manage items)..... 7
 - 7.5 Dashboard page (analytics + reporting)..... 8
 - 7.6 Admin page (user management)..... 8
- 8. V&V Plan..... 8
 - 8.1 Frontend UI Component..... 8
 - 8.2 Backend Services (Inventory Management)..... 9
 - 8.3 Dashboard and Reporting..... 9
 - 8.4 Authentication and Authorization..... 10
 - 8.5 ML Demand Forecasting Component..... 10
 - 8.6 Data Storage..... 11
 - 8.7 Mobile App..... 11
- Appendix..... 13
 - 1. Component Diagram..... 13
 - 2. Component Descriptions..... 13
 - 3. UI/UX Design..... 18

2. Version History

Version	Authors	Change	Date
0	Entire Team	Initial Draft	January 18, 2026
1	Entire Team	Updated Unit Testing Results. In Section 8 under the V&V Plan a table is added for each component with headings <i>Verification Requirement, Test Strategy / Scenarios, Status</i> , and	April 7, 2026

		<i>Result Notes.</i>	
--	--	----------------------	--

3. Purpose Statement

The purpose of this project is to develop an integrated inventory management system that streamlines resource tracking for non-profit organizations with a focus on improving efficiency and accessibility. The system provides a web-based dashboard for managing inventory, visualizing summary analytics, and supporting data-driven decision making through machine learning-driven demand forecasting, alongside a mobile application that enables barcode scanning using device cameras. By replacing manual record keeping and costly proprietary software with a free, open-source solution, the project aims to reduce operational overhead for non-profits so these organizations can spend more time on community impact. The purpose of this documentation is to provide a detailed technical description of the systems components, their relationships between each other and requirements, as well as their underlying implementations in the current codebase. Additionally, this document outlines a Validation and Verification (V&V) plan, detailing the unit testing methodologies and performance metrics to ensure that the software meets all the functional and non-functional requirements defined in the Software Requirements Specification (SRS) document.

4. Component Diagram

See image in appendix section 1.

5. Relationship Between Components and Requirements

5.1. Backend

Requirement	Related Component(s)	Description
Inventory Management - P0	API : Inventory, Database : inventory, locations	Implemented via <code>/inventory/add</code> , <code>/inventory/{item_id}</code> (GET/PUT/DELETE), <code>/inventory/all</code> , <code>/inventory/allwithlocation</code> endpoints; inventory data is stored in the <code>inventory</code> table with foreign key to <code>locations</code> .
Barcode Lookup & Registration - P0	API : Barcode, Database : inventory	Implemented via <code>/barcode/{barcode}</code> GET endpoint barcode field stored with inventory items to support lookup and reuse.
Barcode Scanning (In / Out) - P0	API : Barcode, Database : inventory, inventory_movement	Two step scan-in and scan-out flows: Scan in: 1. <code>/barcode/scan-in</code> (lookup) 2. <code>/barcode/{item_id}/increase</code> (confirmation) Scan out: 1. <code>/barcode/scan-out</code> (lookup) 2. <code>/barcode/scan-out/{item_id}/confirm</code> (confirmation) Updates <code>inventory</code> quantities and creates <code>inventory_movement</code> logging records.
Logging Mechanism / Traceability - P0	Database : inventory, inventory_movement	Every quantity change (scan-in, scan-out, adjustment) is recorded as an immutable entry either in <code>inventory</code> or <code>inventory_movement</code> .
Role-Based Access / Authentication - P1	API : Auth Database : users, roles, food_banks	User identity and role information stored in <code>users</code> and <code>roles</code> ; endpoints are protected to ensure only authorized users can modify inventory.

Demand Forecasting, Procurement Advising - P1	API: Inventory, Database : inventory_movement, inventory	Historical inflow/outflow data required by forecasting models is derived from inventory movement logs and current stock levels. Inventory quantities provide the “current stock” input used by forecasting and procurement logic.
Audit & Performance Reporting - P2	Database : inventory_movement	Inventory movement logs already support auditability; additional aggregation endpoints can be layered on for yearly summaries.
Multi-location Support - P2	Database : locations, API : Inventory	Inventory items already reference locations; supports future expansion to multiple warehouses per food bank.
Performance Considerations - P2	API : Inventory, Database	Separation of inventory and movement tables enables efficient querying and indexing for future scaling.
Mobile App - P2	API : Inventory, Barcode, Auth, Database : inventory, , inventory_movement, locations, users, roles, food_banks	The mobile app acts as a thin client that sends scanned barcode data to the backend APIs, triggering inventory updates, validation, and inventory movement logging in the same way as the web platform.

5.2. Frontend

Requirement	Related Component(s)	Description
Authentication & Role-Based Access - P1	Login Page, Admin Page (Authentication)	Login page for employees, ensures only authorized persons can access/manage inventory for a location associated with the user (accessed via foreign key in user database to index into food_banks). The admin page allows admins to view and manage users and their permissions.
Inventory Management- P0	Inventory Page	Dedicated page for inventory access and management, displays current inventory in table format. Users can add, delete, or edit items and filter tables by name, category, stock level, expiration date, etc.
Barcode Scanning UI - P0	Barcode Scanner	Users can interact with the plus (+) on the home page to enter the scanning interface to scan items in, out, and bulk upload (new items enter the add item window).
Dashboard - P0	Dashboard Page	Page displaying overviews of stock levels and category distribution via charts (with options for filtering) and low stock & expiring items via tables with a slider to set a stock level/expiration date threshold. “Download PDF.” button allows user to export inventory data
Demand Forecasting - P1	Dashboard Page	The user interacts with charts, and gets data and insights generated by the demand forecasting model on the dashboard page.

Mobile App - P2	Barcode Scanner	The mobile page provides a simplified interface for barcode scanning and inventory actions, allowing users to quickly scan items and confirm updates with minimal interaction.
-----------------	-----------------	--

6. Component Descriptions

6.1. Backend Components

The below table summarises the backend component as a whole. Further details about each subcomponent within the backend can be found in section 2 of the Appendix.

Aspect	Description
Normal Behavior	The backend component handles all server-side logic for inventory management, barcode-based workflows, authentication, and data persistence. It validates incoming requests, enforces role-based access control and food-bank scoping, updates inventory state, records inventory movements for traceability, and returns structured responses to the frontend.
API (Inputs / Outputs)	The backend exposes RESTful endpoints for inventory management (<code>/inventory/*</code>), barcode operations (<code>/barcode/*</code>), authentication, and inventory analytics. Inputs are structured JSON payloads containing item metadata, quantities, barcodes, user context, and location identifiers. Outputs are structured inventory and lookup objects used by the frontend to render inventory lists, barcode scan confirmations, and dashboard views.
Implementation	Implemented using a FastAPI-based backend with modular route files (e.g., inventory and barcode routes). Data is persisted in a PostgreSQL database consisting of tables for food banks, users, roles, locations, inventory items, and inventory movement logs. Foreign key relationships enforce organizational scoping and referential integrity. Inventory updates and barcode actions interact with both the inventory and inventory_movement tables to maintain auditability.
Potential Undesired Behaviours	Invalid input data (e.g., negative quantities, malformed barcodes) may cause inconsistent inventory states if not validated. Concurrent inventory updates may lead to race conditions or incorrect quantities. Missing or inconsistent movement logging may reduce auditability and affect forecasting accuracy. Failure to enforce food-bank scoping or role checks could lead to unauthorized data access. Performance issues may arise as inventory and movement records grow without proper indexing.

6.2. Frontend Components

Aspect	Description
Normal Behavior	The frontend component will handle all client-side presentation and interactions for displaying and managing up-to-date food bank inventory. It provides an interface for updating inventory (scanning barcodes, editing current inventory), and provides an additional admin page specific to users with an administrative role. The mobile app frontend provides a streamlined scanning-first workflow where users authenticate, scan an item in/out, review the pre-filled item details (if available), confirm the update, and immediately see the updated quantity reflected in-app and in the web dashboard.
API / Interfaces (Inputs & Outputs)	The frontend communicates with the backend via REST endpoints in order to manage inventory, authenticate users, and display analytics. Inputs consist of JSON payloads and query parameters containing item metadata, quantities, barcode values, movement context, filters, and implicit user and food-bank scope. Outputs are structured JSON responses containing inventory records, barcode scan

	results, inventory movement confirmations, analytics, and error or validation messages, which the frontend uses to update views and provide user feedback. The mobile app exchanges a small set of JSON objects with the backend: it sends <code>{barcode, action, quantity, location_id?}</code> (plus optional item fields for new items) and receives either a “candidate item” object for user confirmation or a finalized <code>InventoryRead</code>/confirmation object (plus validation/error messages) that drives the UI state.
Implementation	Implemented as a client-side application with components for inventory lists, barcode scanning workflows, and dashboards. API communication is handled through a centralized service layer responsible for request construction, authentication headers, and response handling. State management tracks inventory data, loading states, errors, and user context. UI components consume normalized response data to ensure consistent rendering and interaction patterns across views. Implementation is organized around a few screen-level components (e.g., <code>ScannerScreen</code>, <code>ConfirmModal</code>, <code>ResultView</code>) that call thin client service classes (e.g., <code>BarcodeApiClient</code> and <code>InventoryApiClient</code>), where the scanner produces a normalized barcode string, the barcode client fetches candidate/known-item data, and the inventory client performs the confirmed quantity adjustment using the same shared auth/session state for request headers.
Potential Undesired Behaviors	Invalid or incomplete frontend inputs (e.g., missing required fields, incorrect quantities, malformed barcode values) may result in backend validation errors or failed operations. Inadequate handling of concurrent updates or stale state may cause the UI to display outdated inventory quantities. Poor error handling or unclear feedback may confuse users during barcode scans or inventory adjustments. Performance issues may arise if large inventory datasets or analytics responses are rendered without pagination, filtering, or efficient state updates. Potential undesired behaviors include camera/permission failures, unreliable scans producing malformed/partial barcodes, duplicate rapid submissions causing double-counting, stale local state showing outdated quantities under concurrent updates, and network drops leading to lost or silently-unapplied updates if retries/confirmations aren’t handled.

6.3 Machine Learning Component

Aspect	Description
Normal Behavior	When the dashboard loads, the system checks whether a forecast run exists for the food bank within the last 7 days. If the forecast is stale or absent, a background task retrains the model automatically without blocking the UI. The chart renders historical weekly outbound distribution totals alongside point forecasts and confidence intervals for the next 8 weeks, broken down by category. Under normal operation, the model retrains silently, and the chart updates on the next page load.
API (Inputs / Outputs)	GET /forecasts/category?bank_id=1&category=Canned+Goods returns the most recent forecast run for that category: historical weekly totals, point forecasts, and 80%/95% confidence intervals as dated series. POST /forecasts/run manually triggers a full retrain for all categories. Both endpoints are authenticated. Inputs are the bank ID and an optional category filter. Outputs are structured JSON containing forecast values, model metadata (smoothing parameters, training window), backtest WAPE, and a model health indicator (good, degraded, poor).

Implementation	The model is Holt-Winters Exponential Smoothing (ETS A,A,A), additive error, trend, and 52-period weekly seasonality, fit per category using Maximum Likelihood Estimation via the statsmodels library. Parameters α (level), β (trend), and γ (seasonal) are optimised automatically on each food bank's own outbound movement history. The system falls back to non-seasonal ETS when fewer than 104 weeks of data exist, and to a 4-week rolling mean if fitting fails entirely. Results are persisted to the forecast_runs and forecast_values tables and served from cache on subsequent requests. In the event of sparse data, the select_forecaster method will switch to a Croston TSB model, which iteratively tracks the probability of a demand occurring and the expected size of demand, and outputs the product of these two for a stable, low-level estimate.
Potential Undesired Behaviours	If a food bank has fewer than two years of data, the seasonal component cannot be learned, so the model produces a flat or trend-only forecast that will miss annual demand spikes. Sudden one-time events, large donations, and emergencies can distort the smoothed level and produce elevated forecasts for several weeks afterward. Categories with highly intermittent demand (many zero weeks) may receive poor ETS fits; the system detects this and switches to Croston TSB, but accuracy remains low on sparse series regardless of model choice.

7. User Interface

Our UI is designed to be clean and quick to use for day-to-day warehouse work, with the main goal being: users should immediately see what needs attention, then complete common actions (scan in, scan out, add/edit items, view trends) with minimal clicks.

7.1 Layout, navigation, and visual style (*Figures 3.1, 3.1.2 in Appendix*)

Global navigation

- Fixed top navigation bar across the application with tabs: **Home**, **Inventory**, **Dashboard**, and **Admin** (admin-only).
- User account icon on the top-right provides a dropdown for **Profile** and **Logout**.

Design system

- Neutral background with white card containers, rounded corners, and soft shadows for separation.
- Dark navy/charcoal header and primary buttons to create hierarchy.
- Accent colours such as orange and bright purple used mainly for data and status cues (ex: low stock warning icon, expiring soon highlight, chart colours).
- Consistent sans-serif typography with clear hierarchy: large page titles, medium card headers, and smaller table/form text.

Primary actions

- Floating “+” button on the bottom-right provides fast access to the most common workflows: **Scan In** (Add/Update), **Scan Out**, and **Bulk Import**.

7.2 Authentication and account actions (*Figures 3.2, 3.3, 3.4 in Appendix*)

Sign-in page

- A centered “Welcome Back” card with two required fields (Email, Password) and a single primary action button (**Sign In**).
- Layout is intentionally minimal to reduce friction and keep focus on authentication.

Profile and security

- The Profile page (“My Profile”) displays account information (Name, Role, Bank ID) in a simple card layout.
- A Security card includes **Change Password**, which opens a modal requiring current password, new password, and confirmation, with a minimum-length note and clear Cancel/Update actions.

7.3 Home page (overview + forecasting) (*Figures 3.1, 3.5, 3.6, 3.7 in Appendix*)

Top snapshot

- KPI cards at the top: Total Items, Low Stock, Expiring Soon, and This Month Distributed.

- These are placed first since they are the highest priority and most actionable.

Operational cards

- Category Distribution card (left) and Recent Inventory Activity table (right), allowing users to quickly understand current stock makeup and recent updates.

Forecasting & Trends

- Lower section contains forecasting graphs with category dropdown selectors.
- Placed below the operational snapshot because it supports planning, while the KPI cards support immediate action.

Scan In flow (from the “+” button)

- Selecting **Scan In** opens a scan overlay with a camera preview, a detected/manual barcode field, and clear actions.
- After “Use this code,” a confirmation modal appears with the matched item information and editable fields (quantity, unit, name, category, expiry) so users can correct database mismatches before confirming.

7.4 Inventory page (search + manage items) (Figures 3.8, 3.9 in Appendix)

Inventory table

- Table-based layout with columns for item identity, quantity/unit, expiry, location and metadata (date added/modified).
- Edit and delete actions appear as icons per row to keep the table compact and scannable.

Filters

- Search bar and filter controls live above the table (category, low stock threshold, expiry filters) so users can narrow results before scanning rows.

Add item

- “Add Item” button in the top-right opens a modal form with required fields, dropdowns (category) and a date picker (expires), with clear Cancel/Add actions.

7.5 Dashboard page (analytics + reporting) (Figure 3.10 in Appendix)

High-level analytics

- Two primary charts at the top (stock levels and category distribution) with category filtering.
- **Download PDF** button placed top-right for reporting/export.

Attention lists

- Lower cards highlight Low Stock and Expiring Soon items.
- Sliders allow adjusting thresholds and time windows without changing system settings.

7.6 Admin page (user management) (Figures 3.11, 3.12 in Appendix)

Purpose

- Provides administrative control over system access, primarily through user creation and status visibility.

User Management table

- Table lists users with columns: Name, Email, Role, Bank ID and Status.
- Status is shown as a clear badge (ex: “Active”) to make account state visible at a glance.

Create User

- “Create User” button opens a modal form with required fields (Name, Email), Bank ID and a Role dropdown.
- Modal uses the same pattern as other forms: centered card, dimmed background, and clear Cancel/Primary action (“Create User”) buttons for consistency.

Mobile (early prototype) (Figure 3.13 in Appendix)

- In addition to the web UI, we have started an early mobile prototype focused on the highest-frequency task: barcode scanning. The current app screen supports **Scan In (Add)** and **Scan Out (Remove)** with a live camera view and a simple loading/processing state while a barcode is being detected. As development continues, the mobile app will follow the same visual design system as the web application (navigation style, colours, typography, and component patterns) to keep the experience consistent across platforms.

8. V&V Plan

8.1 Frontend UI Component

This is the main website volunteers use to manage items. We need to ensure it is intuitive and error-free as defined in SRS Section 9.1.

Unit Tests (Logic & Behavior) Tools: Jest and React Testing Library

- **Item Entry and Editing Forms**
 - Verify that required fields (item name, category, quantity) must be filled before submission.
 - Ensure numeric fields reject invalid input:
 - Negative quantities
 - Non-numeric values
 - Confirm that expiration dates cannot be set in the past.
 - Validate that category selections are correctly stored and reflected in the UI.
 - Test that successful form submission triggers a confirmation state or success message.
- **Barcode Scanning Interface**
 - Simulate barcode scan events and verify that recognized barcodes automatically populate item fields (name, category, unit if available).
 - Confirm that immediate visual feedback is displayed upon a successful scan (e.g., green checkmark or confirmation indicator), in line with the **Look and Feel requirements** in SRS Section 11.
 - Test fallback behavior when a barcode is not recognized:
 - Manual entry fields become available
 - Clear messaging is shown to guide the user
- **Error Handling and User Feedback**
 - Ensure validation errors are displayed clearly and in plain language (e.g., “Quantity must be greater than zero”); also, that users are clearly notified of authentication errors.
 - Verify that error states do not clear previously entered valid data.
 - Confirm that users can correct mistakes without restarting the workflow, supporting the “forgiving to mistakes” usability requirement.
- **Button and Interaction States**
 - Verify that submit buttons are disabled until all required fields are valid.
 - Test loading or disabled states while actions are in progress to prevent duplicate submissions.
 - Ensure cancel or reset actions behave as expected.

Performance Metrics (Speed)

Tools: Chrome DevTools


Metric: Time to Interact

Goal: ≤ 1.5 seconds

The time required for the page to become fully interactive will be measured using Chrome DevTools performance profiling under simulated low-performance conditions, ensuring usability on older hardware commonly used in food bank environments.

Verification Requirement	Test Strategy / Scenarios	Status	Result Notes

Item Entry & Validation	Rendered <code><AddItemModal></code> and utilized <code>userEvent</code> to attempt submissions with invalid string inputs in numeric fields (<code><input type="number" min="0"></code>). <code>/frontend/src/__tests__/AddItemModal.test.jsx</code>	Passed	Native HTML constraints successfully trapped invalid data. Negative values and non-numeric strings are safely rejected (cleared to null) without crashing the component.
Error Handling States	Simulated an "Edit" mode scenario that violates business logic (decreasing quantity while simultaneously transferring location). <code>/frontend/src/__tests__/AddItemModal.test.jsx</code>	Passed	The system successfully halted form submission and rendered an explicit warning banner above the form ("Cannot change both Quantity and Location"). Valid data was preserved, allowing the user to recover without restarting.
Barcode Scan Population	Simulated a 12-digit barcode input stroke-by-stroke. Mocked the <code>fetchInventoryByBarcode</code> service to return a registered database item. <code>/frontend/src/__tests__/AddItemModal.test.jsx</code>	Passed	Verified that the API resolves successfully. State dynamically updates to auto-populate "Name" ("Heinz Baked Beans") and paints the visual confirmation flag ("Barcode already in inventory").
Interaction States	Rendered the modal with the <code>isSaving={true}</code> prop, simulating an active API network request. <code>/frontend/src/__tests__/AddItemModal.test.jsx</code>	Passed	The Submit button dynamically updated to "Saving..." and both Submit and Cancel buttons were immediately set to disabled. This safely mitigates the risk of duplicate database entries or aborted network requests.

Performance Metrics (Speed): Time to Interact (TTI)	Utilized Chrome DevTools Performance Profiler with CPU Throttling set to 4x Slowdown to simulate a low-spec device. Recorded the initial mount and DOM render lifecycle.	Passed	Despite CPU throttling, the React application parsed and resolved the main thread under the threshold. TTI was recorded at 1.374 seconds, successfully under the 1.5 seconds requirement. 
--	---	---------------	--

8.2 Backend Services (Inventory Management)

The inventory management backend handles all core data processing, including inventory updates, calculations, and bulk uploads.

Unit Tests (Logic & Behavior) Tools: PyTest

- Calculation Accuracy**
 - Inventory totals update correctly when items are added or removed.
 - No incorrect arithmetic occurs during repeated updates.
- Bulk Upload Validation**
 - Correctly formatted CSV/Excel files are processed successfully.
 - Malformed files (missing columns, invalid values) are rejected safely.
 - Partial uploads do not corrupt existing inventory data.

Performance Metrics (Responsiveness)

Tools: Scripts & Postman


Metric: API Response Time.

Goal: When a volunteer scans an item, the backend should send back the item details in under **200 milliseconds**. This ensures the scanning process feels instant.

Performance validation will focus on API response time, measured using simple request scripts and Postman for manual verification.

Verification Requirement	Test Strategy / Scenarios	Status	Result Notes

Calculation Accuracy: Add/Remove	Simulated INBOUND (+ delta) and OUTBOUND (- delta) adjustments via <code>adjust_item_quantity</code>. <code>/backend/tests/test_inventory_service.py</code>	Passed	Inventory totals successfully update and correctly reflect positive and negative deltas. Strict bounds checking prevents negative stock anomalies (throws HTTP 400 Insufficient stock).
Calculation Accuracy: Repeated Updates	Programmatically executed rapid, sequential addition and subtraction operations on a single <code>item_id</code>. <code>/backend/tests/test_inventory_service.py</code>	Passed	Arithmetic integrity holds under repeated updates. Validated that sequential mutations (-5, +20, -30) compute safely with no arithmetic drifting.
Movement Logging Auditability	Queried the <code>inventory_movement</code> table after repeated mathematical adjustments. <code>/backend/tests/test_inventory_service.py</code>	Passed	Verified that every arithmetic operation triggers a corresponding immutable entry in the movement log, maintaining 1:1 traceability
Bulk Upload Validation: Valid CSV Processing	Supplied a correctly formatted byte-stream mimicking a CSV upload to the parser to verify data mapping. <code>/backend/tests/test_bulk_import.py</code>	Passed	100% of rows were successfully parsed, mapped to the <code>InventoryCreate</code> schema, committed to the database, and returned a success count matching the <code>total_items</code> .
Bulk Upload Validation: Malformed File Rejection	Simulated a file upload containing strict integer columns populated with string values (e.g., <code>quantity="not_a_number"</code>) and missing fields.	Passed	The system successfully intercepted the parsing exceptions (catching <code>ValueError</code> and <code>KeyError</code>). The endpoint returned a detailed error manifest and zero corrupt rows were committed to the database.

	/backend/tests/test_bulk_import.py		
Bulk Upload Validation: Partial Upload Integrity	Simulated a mixed file upload containing both valid records and explicitly corrupt records. /backend/tests/test_bulk_import.py	Passed	The loop-level transaction handling successfully committed the valid rows while catching and skipping the invalid rows. Verified that a single corrupt row does not roll back the entire valid batch, preserving partial import capabilities.
Performance Metrics: API Response Time	Execute an automated Python request script utilizing a valid JWT authentication token to simulate a volunteer rapidly scanning 100 sequential items under isolated, local conditions to capture baseline network latency and ensure resolution is well under the 200ms threshold. /backend/test_performance.py	Passed	— Performance Results — Target Goal: < 200.00 ms Average Response Time: 7.89 ms Min Response Time: 4.77 ms Max Response Time: 44.30 ms 

8.3 Dashboard and Reporting

The dashboard and reporting component presents aggregated inventory data to support staff decision-making. Verification ensures that displayed information is accurate and consistent with stored data.

Unit Tests (Logic & Behavior) Tools: PyTest

- **Data Aggregation**
 - We will verify that the charts (e.g., "Total Inventory by Category") accurately sum up the data from the database and don't miss any records. This will be done by providing controlled input datasets and checking that computed summaries, such as total inventory by category, match expected results.
- **Alert Triggers**
 - We will test that "Low Stock" and "Expiring Soon" alerts actually appear on the dashboard when items cross the defined thresholds (e.g., expiring within 30 days).

Performance Metrics (Load Time)

Tools: Scripts

Metric 1: Report Generation Speed, defined as the time required to generate and prepare reports for download. This will be measured by recording timestamps before and after report generation in the backend.

Goal: Complex reports (like the End-of-Year Audit) should generate and be ready for download within **5 seconds**, so staff members aren't left waiting.

Metric 2: Chart Rendering Time

Goal: ≤ 500 milliseconds per chart. This ensures that charts appear quickly and do not feel “laggy” or unresponsive. Chart rendering time will be measured by:

- Logging timestamps before and after chart rendering in the frontend, or
- Using Chrome DevTools to inspect component render timing.

Verification Requirement	Test Strategy / Scenarios	Status	Result Notes
Data Aggregation	Seeded the database with specific item categories and called <code>get_category_group_summary</code> to verify output sums. <code>/backend/tests/test_dashboard_reporting.py</code>	Passed	Assertions confirmed the endpoint sums categorized inventory without skipping or duplicating records (e.g., 20 Apples + 25 Bananas aggregated safely to exactly 45 Produce).
Alert Triggers (Thresholds)	Verified boundary limits by generating mock items with specific quantities and expiration dates relative to <code>date.today()</code> . <code>/backend/tests/test_dashboard_reporting.py</code>	Passed	Database queries confirmed strict adherence to thresholds. "Low Stock" correctly isolated only items <code><= 10</code> units, and "Expiring Soon" isolated only items expiring within <code>30</code> days.
Performance Metric: Report Generation Speed	Automated script executed a massive "End-of-Year Audit" query (<code>/reports/download</code> for a 2-year window), recording timestamps directly before and after the network request. <code>/backend/test_report_performance.py</code>	Passed	Report compiled, converted to CSV, and returned successfully in 0.04 seconds , surpassing the <code><= 5</code> -second complex report goal.

Performance Metric: Chart Rendering Time	Manual Verification: Utilized Google Chrome DevTools (Performance Profile Tab). Navigated to the Dashboard and recorded the render timeline for the Recharts components.	Passed	DevTools profiling confirmed the DOM paints the "Total Inventory by Category" and "Demand Line Chart" SVGs in under ~ 150ms , well below the 500ms threshold.
---	---	---------------	--

8.4 Authentication and Authorization

The authentication component ensures that users can only access features appropriate to their roles, protecting sensitive inventory data.

Unit Tests (Logic & Behavior) Tools: PyTest

- **Access Control:**
 - We will ensure that *Volunteer* users are unable to access *Admin* settings, that unauthorized users are completely blocked from viewing inventory data, and that authorized users can view the inventory information relevant to their role.
- **Login Security:**
 - We will test that the system correctly rejects invalid passwords and securely manages user sessions (using tokens) so users don't have to log in constantly.

Performance Metrics (Security Latency)

Tools: Scripts

Metric: Login Handshake Time, measured by timing the authentication process across multiple login attempts.

Goal: The security checks during login should add less than **100 milliseconds** to the login process, keeping security from becoming an annoyance.

Verification Requirement	Test Strategy / Scenarios	Status	Result Notes
Access Control (Unauthorized Bypass)	Attempted to submit <code>POST</code> requests to protected routes (e.g., <code>/barcode/scan-out</code>) without appending a JWT Bearer token to the header. <code>/backend/tests/test_auth_authorization.py</code>	Passed	The API dependency injection layer successfully intercepted the request prior to execution, returning an HTTP <code>401 Unauthorized</code> block. Inventory data remained entirely inaccessible to unauthenticated traffic.

Login Security (Invalid Credentials)	Submitted an OAuth2 form payload containing a registered email but an explicitly incorrect string for the password field. /backend/tests/test_auth_authorization.py	Passed	The verify_password hashing check successfully detected the mismatch and rejected the authentication attempt, returning a 401 error flag. No token was issued.
Session Management (Token Generation)	Submitted valid admin credentials to /auth/login and subsequently passed the returned payload to a secondary protected route. /backend/tests/test_auth_authorization.py	Passed	The backend securely authorized the user and generated a standard OAuth2 JWT Access Token. The token successfully bypassed secondary route security checks, validating the session management flow.
Performance Metrics (Security Latency): Login Handshake Time	Executed an automated Python script to perform 50 sequential, successful login handshakes against the /auth/login endpoint, tracking the full network round-trip and database query lifecycle. /backend/test_auth_performance.py	Failed (Accepted Variance)	Target Goal: < 100.00 ms Average Handshake: 212.50 ms Min Handshake: 207.21 ms Max Handshake: 383.09 ms 95th Percentile: 210.34 ms Engineering Justification: The system slightly exceeds the 100ms target, but this is a deliberate and necessary trade-off for security. We use the bcrypt hashing algorithm with a standard cost factor (rounds = 12), which intentionally adds around 200ms of processing time to protect against brute-force attacks. Reducing login time below 100ms would mean lowering this cost factor, which would weaken the system's security. Since a response time of about 212ms still feels instant to users, we've chosen to prioritize strong security and data

			protection over meeting a strict performance threshold.
--	--	--	---

8.5 ML Demand Forecasting Component

The demand forecasting component predicts future inventory needs based on historical data. Verification focuses on ensuring that the data pipeline feeding the model is reliable and well-structured.

Unit Tests (Logic & Behavior) Tools: PyTest

- **Data Cleaning and Formatting**
 - Historical inventory data is correctly parsed from storage into the expected format.
 - Missing values (e.g., missing dates or quantities) are handled safely through removal or imputation.
 - Invalid records (negative demand, malformed timestamps) are rejected or corrected.
 - Item categories are consistently encoded across time periods.
- **Feature Engineering Validation**
 - Date-based features (e.g., month, week, season) are generated correctly.
 - Holiday and seasonal indicators are applied consistently to historical records.
 - Perishable and non-perishable items are handled according to their defined characteristics.
- **Pipeline Consistency**
 - The same preprocessing steps are applied to both training data and future prediction inputs.
 - Input feature dimensions remain consistent across runs.
 - The pipeline produces deterministic outputs when given the same input data.

Performance Metrics (Accuracy)

Tools: Scripts

Metric 1: Weighted Absolute Percentage Error (WAPE)

Goal: < 30%

WAPE measures the total absolute prediction error as a percentage of total actual demand. This metric is well-suited for inventory forecasting because it accounts for scale and provides an intuitive sense of overall prediction quality. A WAPE below 30% indicates that forecasts are reasonably close to actual demand and suitable for decision support.

Metric 2: Forecast Bias

Goal: within $\pm 10\%$


Forecast bias measures whether the model systematically overestimates or underestimates demand. Keeping bias within $\pm 10\%$ ensures the system does not consistently recommend overstocking (leading to waste) or understocking (leading to shortages).

Performance metrics will be computed using simple Python evaluation scripts that compare model predictions against known historical demand values.

Metric 3: Mean Absolute Error (MAE)

Goal: ≤ 3 units per item

Mean Absolute Error (MAE) measures the average absolute difference between predicted demand and actual demand, expressed in the same units as the inventory items. A target MAE of 3 units or less indicates that forecasts are typically within a small, operationally acceptable margin of error, making them suitable for short-term procurement and inventory planning decisions.

Verification Requirement	Test Strategy / Scenarios	Status	Result Notes
ML Performance Metrics (Accuracy)	expanding-window backtest on 110 weeks of real McMaster Foodbank data, evaluating CrostonTSB at h=1, h=2, and h=4 week horizons; results: MAE=77.3 units/week, WAPE=87%, Bias=-16%.	Failed	Because the McMaster food bank data we procured was too sparse, we were only able to evaluate the Croston and naive average models (not the Holt-Winters ETS with seasonality). In the final production environment, the ETS model is used based on the food bank's own data, after 2 years (cycles) annual seasonality is visible. With synthesized data that has annual spikes, seasonality is capture with metrics: MAE 19.0 units/week, WAPE 9.9%, MASE 0.89. Although synthesized, these results show how, in production, seasonality is captured by ETS.
Data Cleaning & Formatting	Executed PyTest suite (/backend/tests/forecasting/test_pipeline.py) injecting malformed pandas DataFrames containing missing timestamps, null quantities, and invalid negative outbound demand.	Passed	The preprocessing pipeline successfully dropped rows with unresolvable missing dates and strictly filtered out invalid negative demand values. Categorical variables (e.g., item categories) were verified to encode identically across distinct historical data batches. 
Feature Engineering Validation	Evaluated the timestamp extraction functions and category-mapping logic in isolation using PyTest. Verified the generation of derived time-series columns (month, week). /backend/tests/forecasting/test_pipeline.py	Passed	Date-time components were accurately extracted into independent numerical features (Month, Day of Week). The pipeline correctly mapped categories (e.g., Perishable Produce vs. Canned Goods) to ensure the forecasting model can detect shelf-life and seasonal demand spikes.

Pipeline Consistency	Processed identical datasets through both the historical training data path and the future prediction data path. Asserted that both resulting feature matrices possessed identical dimensionality. /backend/tests/forecasting/test_pipeline.py	Passed	Both execution paths yield exactly matching feature spaces. The pipeline guarantees that real-time prediction inputs undergo the exact same scalar transformations as the historical training data. Running the same data twice yielded 100% deterministic outputs.
-----------------------------	--	---------------	--

8.6 Data Storage

The data storage component manages all persistent system data, including inventory records and historical logs. Verification ensures that data remains consistent under concurrent access.

Unit Tests (Logic & Behavior) Tools: PyTest

- **Integrity Checks**
 - We will test that data remains consistent even if two users try to update the same item at the exact same time.
- **Backup/Recovery**
 - We will verify that if we delete data appropriately, it is gone, but if we need to query historical logs, they are still available.

Performance Metrics (Capacity)

Tools: Scripts to issue multiple concurrent queries

Metric: Query Latency under Load.

Goal: The database should handle 50 concurrent requests (simulating multiple busy volunteers) without slowing down significantly.

Verification Requirement	Test Strategy / Scenarios	Status	Result Notes
Integrity Checks (Concurrency)	Executed rapid, sequential arithmetic updates (+5, -2, +10) on the <code>InventoryItem.quantity</code> column within the same database engine scope. /backend/tests/test_data_storage.py	Passed	SQLAlchemy transaction isolation successfully maintained arithmetic integrity. The final quantity resolved to the exact expected sum (23), proving no "lost updates" occurred during rapid commits.

Backup/Recovery (Audit Logs)	Programmatically committed a transaction to the <code>inventory_movement</code> table and verified historical persistence. <code>/backend/tests/test_data_storage.py</code>	Passed	Verified that operations successfully generate entries in the <code>inventory_movement</code> log with correct <code>quantity_change</code> and <code>movement_type</code> tags. The database structurally retains these logs for audit trails and recovery.
Performance Metrics (Capacity) : Query Latency under Load	Used <code>ThreadPoolExecutor</code> to launch 50 parallel asynchronous <code>GET</code> queries against the inventory list endpoint. <code>/backend/test_db_load.py</code>	Passed	<pre> Starting Database Capacity Test... URL being tested: http://127.0.0.1:8000/inventory/all Simulating 50 simultaneous database queries... --- Database Load Results --- Total Requests: 50 Successful Resolves: 50/50 Target Goal: < 500 ms p95 Average Latency: 127.15 ms Max Latency: 144.49 ms 95th Percentile: 143.44 ms </pre> The system handled 50 concurrent requests without connection pooling errors or significant performance degradation, maintaining sub-100ms response times.

8.7 Mobile App

This is the portable version of the tool for scanning on the go.

Unit Tests (Logic & Behavior) Tools: Mobile emulators, mocked camera input, barcode scanning libraries

- **Camera Access**
 - The app successfully requests and receives camera permissions.
 - Clear error messaging is displayed if camera access is denied.
 - The app recovers correctly when permissions are later granted.
- **Barcode Recognition**
 - Barcodes are correctly recognized using mocked camera frames.
 - Recognition remains reliable under:
 - Low or uneven lighting
 - Slight camera movement
 - Different barcode orientations
 - Successful scans trigger immediate visual feedback (e.g., confirmation indicator).
- **Manual Fallback**
 - If a barcode cannot be recognized, the app allows manual item entry.
 - Manual entries follow the same validation rules as the web interface.
- **Real-Time Sync**
 - Items scanned on the mobile app appear on the web dashboard immediately after submission.
 - Quantity changes made on mobile update aggregate counts correctly.
- **Consistency Across Clients**

- Inventory state remains consistent when multiple devices update data concurrently.
- Duplicate submissions are prevented if the same item is scanned multiple times rapidly.
- **Network Disruption Handling**
 - If connectivity is temporarily lost, the app:
 - Clearly notifies the user
 - Prevents silent data loss
 - Once connectivity is restored, pending updates are either safely retried or require user confirmation.

Performance Metrics (Startup)

Tools: Scripts & Built-in development environment tools

Metric 1: App Launch Time

Goal: ≤ 2 seconds, supporting efficient use during high-activity intake periods.

Metric 2: Scan-to-Submit Latency, the time between a successful barcode scan and confirmation of item submission.

Goal: ≤ 300 milliseconds, ensures the scanning experience feels responsive and does not slow down volunteer workflows.


Verification Requirement	Test Strategy / Scenarios	Status	Result Notes
Camera Access & Permissions	Automated UI test execution using <code>tests/scanner.test.tsx</code> mocking the <code>expo-camera</code> permission lifecycle.	Passed	Verified through Jest that the <code>ScannerScreen</code> correctly branches its UI state. When permissions are denied, the app renders the request screen; when granted, it successfully mounts the <code>CameraView</code>. <pre> PASS tests/scanner.test.tsx Mobile Scanner UI Logic ✓ renders the scanner view when permissions are granted (180 ms) ✓ shows permission request if camera access is missing (16 ms) Test Suites: 1 passed, 1 total Tests: 2 passed, 2 total Snapshots: 0 total Time: 0.708 s, estimated 2 s </pre>
Barcode Recognition Feedback	Simulated scan event triggering <code>setScanned(true)</code> within the <code>handleBarCodeScanned</code> hook.	Passed	Automated state analysis confirmed that a successful scan immediately updates the internal UI state to freeze the camera preview and prompt the confirmation modal.

Manual Fallback Integration	Code analysis of the fallback button handler in <code>index.tsx</code> .	Passed	Verified that the "Manual Entry" button correctly bypasses the camera requirement and opens the input form, following the same schema validation as the web client.
Real-Time Sync & Persistence	Integration testing between the mobile client and the backend API service.	Passed	Confirmed that <code>services/api.ts</code> correctly passes Bearer tokens and JSON payloads. Updates made on the mobile client reflected in the backend database with sub-200ms synchronization latency.
Duplicate Scan Prevention	Automated resilience test in <code>tests/resilience.test.tsx</code> simulating three rapid <code>onBarcodeScanned</code> events fired in succession on the same barcode within a single render cycle.	Passed	Verified through Jest that the loading state flag in <code>handleBarCodeScanned</code> acts as a gate: only the first scan event triggers an API call and opens the confirmation modal. The subsequent two events are dropped. <code>api.post</code> call count was asserted to equal exactly 1 despite 3 scan firings. <pre> PASS tests/resilience.test.tsx Mobile Resilience: Network & Duplicates ✓ prevents duplicate submissions from rapid multi-scans (391 ms) ✓ prevents silent data loss and notifies user on network failure (1074 ms) Test Suites: 1 passed, 1 total Tests: 2 passed, 2 total Snapshots: 0 total Time: 2.217 s Ran all test suites matching /tests\/resilience.test.tsx/i. </pre>
Network Failure & Silent Data Loss Prevention	Automated resilience test in <code>tests/resilience.test.tsx</code> using <code>mockRejectedValueOnce</code> to simulate a network crash ("airplane mode") during the quantity confirmation submit, after a successful scan-in lookup.	Passed	Verified through Jest that when the confirmation <code>api.post</code> rejects with a Network Error, the quantity modal remains visible on screen. The "Quantity to Add" label was asserted to still be present after the failed submit, confirming the volunteer is not misled into thinking the scan succeeded and no inventory data is silently lost. <pre> PASS tests/resilience.test.tsx Mobile Resilience: Network & Duplicates ✓ prevents duplicate submissions from rapid multi-scans (391 ms) ✓ prevents silent data loss and notifies user on network failure (1074 ms) Test Suites: 1 passed, 1 total Tests: 2 passed, 2 total Snapshots: 0 total Time: 2.217 s Ran all test suites matching /tests\/resilience.test.tsx/i. </pre>

Performance Metrics: Scan-to-Submit Latency	Measured the duration between a barcode being recognized and the API "Success" response.	Passed	Actual: 240ms (Goal: ≤ 300 ms). The lightweight mobile payload and optimized backend routes ensure the scanning process feels instantaneous to the user.
Performance Metrics: App Launch Time	Profiled using Expo DevTools recording Time to First Interactive (TTI), which was then verified using physical device cold-start analysis to measure true TTI	Passed	Actual: 1.45s (Goal: ≤ 2 s). Average cold-start to interactive state measured at 1.45s across 3 trials on test hardware. The app reached a fully interactive state well within the threshold.

Appendix

1. Component Diagram



2. Component Descriptions

For software design purposes, some of the functionality and business logic associated with the API endpoints have been moved into service modules, with inventory, barcodes, and auth each having its own respective service module in the backend.

API – Inventory

Aspect	Description
Normal Behavior	The Inventory API provides endpoints for creating, retrieving, updating, deleting, and listing inventory items. It validates incoming item data (e.g., quantity, expiration date, location), adds valid inventory records to the database and returns structured inventory objects to the frontend. Inventory items are always associated with a location and a food bank to prevent cross-organization access.

API (Inputs / Outputs)	POST <code>/inventory/add</code> accepts an inventory payload and returns the newly created inventory item. PUT <code>/inventory/{item_id}</code> accepts updated item fields and returns the modified inventory record. GET <code>/inventory/{item_id}</code> returns a single inventory item by ID. DELETE <code>/inventory/{item_id}</code> removes an inventory item. GET <code>/inventory/all</code> returns a list of all inventory items. GET <code>/inventory/allwithlocation</code> returns inventory items joined with their associated location metadata. All endpoints exchange structured inventory data objects containing item metadata, quantity, expiration date and location information.
Implementation	Implemented using FastAPI routes defined in <code>inventory_routes.py</code>. Each route interacts with the database through a session dependency and performs CRUD operations directly on inventory records. The inventory routes may also trigger creation of corresponding inventory movement records when quantity changes occur. <code>InventoryCreate</code>, <code>InventoryUpdate</code>, and <code>InventoryRead</code> define the request/response shapes; persistence uses <code>InventoryItem</code> and <code>Location</code>. Endpoints call service methods like <code>add_item(...)</code> and <code>update_item(...)</code>, returning <code>InventoryRead</code>.
Potential Undesired Behaviors	Invalid input data (e.g., negative quantities or invalid expiration dates) may cause inconsistent inventory states if not validated. Concurrent updates to the same inventory item may result in race conditions or lost updates. Deleting inventory items without checking movement history may impact auditability.

API – Barcodes

Aspect	Description
Normal Behavior	The Barcode API supports barcode-driven inventory workflows. When a barcode is scanned, the backend attempts to identify the corresponding inventory item and either (1) returns item details for confirmation, or (2) creates/assists creation of a new item if the barcode is not yet associated with inventory. The Barcode API also supports “scan-out” flows that reduce inventory quantities and records the change in the inventory movement log for traceability. Overall, the expected behavior is: barcode → lookup → confirm action → update inventory → log movement → return updated item state.

API / Interfaces (Inputs & Outputs)	GET <code>/barcode/{barcode}</code>: Takes barcode as input and outputs barcode lookup result (inventory item details if found; otherwise not found/404). POST <code>/barcode/add</code>: takes item payload as input and outputs created inventory item associated with that barcode. POST <code>/barcode/scan-in</code>: takes scan request payload containing barcode and quantity as input, outputs an inventory item record reflecting updated quantity (either a newly created item or an existing item updated). POST <code>/barcode/{item_id}/increase</code>: takes item_id + quantity increment payload as input and outputs updated inventory item with increased quantity. POST <code>/barcode/scan-out</code>: takes scan request payload containing barcode and quantity requested for removal as input and outputs lookup response indicating whether an item exists and whether scan-out can proceed. POST <code>/barcode/scan-out/{item_id}/confirm</code>: takes item_id + confirmation payload containing quantity to remove as input and outputs updated inventory item after quantity deduction; also triggers creation of an inventory movement record.
Implementation	Implemented as a FastAPI router in <code>barcode_routes.py</code>. The Barcode API interacts with the inventory subsystem in two ways: (1) direct lookups using the barcode field, and (2) delegated updates to inventory quantities via inventory service/database session operations. The barcode flow centers on <code>BarcodeInfo</code> object and <code>InventoryRead</code> (when a scan leads to an inventory create/update). The main path is <code>upsert_scanned_item(...)</code>, which normalizes the scanned code, looks up an existing <code>InventoryItem</code> under the current user's organization scope, and if none exists, returns a <code>BarcodeInfo</code> candidate to confirm/create. Quantity changes are funneled through a movement type (<code>MovementType.OUTBOUND</code>), linking barcode scans to consistent inventory mutation logic.
Potential Undesired Behaviors	Barcode not found: Scans may not match any stored inventory item, requiring graceful handling (return not found and allow manual item creation). Duplicate or inconsistent barcodes: If barcode normalization is not consistent (leading zeros, whitespace), the system may create duplicates or fail to match an existing item. Missing movement logging: If quantity updates occur without inserting an <code>inventory_movement</code> record, traceability and forecasting inputs become unreliable. Cross-food-bank leakage: If barcode lookup does not enforce scoping rules (<code>bank_id</code>), users may see or modify inventory belonging to another organization.

API – Authentication

Aspect	Description
Normal Behavior	The Auth API supports the fundamental authorization checks as well as Admin activities like creating a user and updating passwords. At login HTTP requests are made to this API to verify credentials and return a valid token. These endpoints also handle creating and updating users (for admins).

API / Interfaces (Inputs & Outputs)	POST <code>/auth/login/</code>: Accepts login credentials and returns a token response representing the validity of the login request. POST <code>/auth/me/</code>: Accepts a token and returns the current user object POST <code>/auth/change-password</code>: Accepts valid current user's credentials, updates current users' password. Authentication not required for this endpoint. GET <code>/auth/users</code>: Returns a list of users for Admin purposes POST <code>/auth/users</code>: Creates a new user with a temporary password, accepts a user object and returns it back on success. Put <code>/auth/users/{user_id}</code>: Accepts a user id parameter and an update object, performs updates, then returns the new user object.
Implementation	Implemented as a FastAPI router in <code>auth_routes.py</code>. The Routes endpoints interact with the authorization services to verify user information for authentication and retrieve and update user information for admin purposes. For authentication, it relies on built-in Python packages in SQLAlchemy (maintaining session and user objects), FastAPI (HTTP framework), Jose (JWT tokens), bcrypt (for hashing). Authentication uses a <code>TokenResponse</code> object for login (including <code>requires_password_change</code>) and the object <code>UserRead</code> for identity endpoints. The relationship between methods is: <code>login(...)</code> authenticates the user, creates an access token, and returns <code>TokenResponse</code>; <code>me(...)</code> returns <code>UserRead</code> object and derives role information via a role-resolution method. User management is authorization-gated (admin-only vs self-update); <code>update_user(...)</code> enforces admin/self rules, fetches the target user by id, applies allowed field updates, commits, and returns <code>UserRead</code> object.
Potential Undesired Behaviors	With respect to security and authorization the potential undesired behaviours are: - Exploitation of rate limiting – at the moment there are incomplete measures against brute force attacks. - Debug Logging currently exposes sensitive user info – this is still the case for development purposes. - Unauthorized access to certain functionality (admin features) or unauthenticated access to the platform (i.e. any security breach) is a potential behaviour.

Database Components – food_banks table

Aspect	Description
Normal Behavior	The <code>food_banks</code> table stores the set of food bank organizations supported by the system.

API / Interfaces (Inputs & Outputs)	Writes: Food banks are created during system setup or onboarding. Reads: Used to scope all authenticated requests (inventory, locations, users) to the correct organization. Data Structure: { <code>bank_id</code> , <code>name</code> , <code>address</code> }.
Implementation	Implemented as a PostgreSQL table. The table is referenced by <code>inventory.bank_id</code> , <code>locations.bank_id</code> , and <code>users.bank_id</code> via foreign key constraints.
Potential Undesired Behaviors	If API queries do not consistently filter by <code>bank_id</code> , data leakage across organizations may occur. Deleting a food bank without restriction rules may orphan dependent rows.

Database Components – locations table

Aspect	Description
Normal Behavior	The <code>locations</code> table stores storage locations associated with a food bank. Inventory items reference locations to support assignment, filtering and future multi-location support.
API / Interfaces (Inputs & Outputs)	Writes: Created or updated when staff manage storage locations. Reads: Used when assigning inventory items to locations and when listing inventory with location metadata. Data Structure: { <code>location_id</code> , <code>bank_id</code> , <code>name</code> , <code>address</code> , <code>notes</code> }.
Implementation	Implemented as a PostgreSQL table The table is referenced by <code>inventory.location_id</code> and by <code>inventory_movement.from_location_id</code> / <code>to_location_id</code> .
Potential Undesired Behaviors	Deleting a location that is still referenced by inventory or movement records may break referential integrity. If <code>bank_id</code> is not validated at the API level, a location from one food bank could be incorrectly assigned to inventory from another.

Database Components – roles table

Aspect	Description
Normal Behavior	The <code>roles</code> table defines the set of roles used for role-based access control (e.g., volunteer, staff, admin). Each user is assigned a role that determines their permitted actions.
API / Interfaces (Inputs & Outputs)	Writes: Typically populated during initial system setup. Reads: Loaded during authentication and authorization checks. Data Structure: { <code>role_id</code> , <code>name</code> }.
Implementation	Implemented as a PostgreSQL table. The table is referenced by <code>users.role_id</code> .

Potential Undesired Behaviors	If role checks are not consistently enforced in API routes, users may gain unauthorized access. Changes to role definitions without corresponding API updates could lead to inconsistent permission behavior.
--------------------------------------	---

Database Components – users table

Aspect	Description
Normal Behavior	The <code>users</code> table stores user identities, hashed passwords, assigned roles and associated food banks. It supports authentication and ensures that all actions are scoped to the user’s organization.
API / Interfaces (Inputs & Outputs)	Writes: User records are created during registration or admin provisioning. Reads: Used during login and authorization checks. Data Structure: { <code>user_id</code> , <code>name</code> , <code>email</code> , <code>role_id</code> , <code>bank_id</code> } (password hash is never returned).
Implementation	Implemented as a PostgreSQL table. Referenced by <code>inventory.created_by</code> and <code>inventory.modified_by</code> .
Potential Undesired Behaviors	Improper handling of sensitive fields (e.g., returning password hashes) poses security risks. If <code>role_id</code> is missing or invalid, permission checks may behave unpredictably. Failure to scope queries by <code>bank_id</code> can lead to cross-organization access.

Database Components – inventory table

Aspect	Description
Normal Behavior	The <code>inventory</code> table represents the current state of all inventory items.
API / Interfaces (Inputs & Outputs)	Writes: Item creation, updates, deletions, and quantity adjustments from scan-in/scan-out operations. Reads: Inventory listings, item detail views, and dashboard aggregations. Data Structure: { <code>item_id</code> , <code>bank_id</code> , <code>name</code> , <code>category</code> , <code>barcode</code> , <code>quantity</code> , <code>unit</code> , <code>expiration_date</code> , <code>location_id</code> , <code>date_added</code> , <code>last_modified</code> , <code>created_by</code> , <code>modified_by</code> }.
Implementation	Implemented as a PostgreSQL table.
Potential Undesired Behaviors	The global uniqueness constraint on <code>barcode</code> may prevent different food banks from storing the same product. Quantity values may become negative if not validated at the API layer. Deleting items that have associated movement records may reduce auditability unless soft-deletes or restrictions are applied.

Database Components – inventory_movement table

Aspect	Description
--------	-------------

Normal Behavior	The <code>inventory_movement</code> table stores an append-only log of all inventory quantity changes. Each record captures how, when, and why an item’s quantity changed, supporting auditability and forecasting.
API / Interfaces (Inputs & Outputs)	Writes: Created whenever inventory quantity changes due to scan-in, scan-out, waste, or transfers. Reads: Used for audit reports, dashboard analytics and demand forecasting inputs. Data Structure: <code>{id, item_id, quantity_change, movement_type, reason, from_location_id, to_location_id, created_at}</code> .
Implementation	Implemented as a PostgreSQL table. Indexed on <code>item_id</code> for efficient lookups.
Potential Undesired Behaviors	Inventory updates and movement inserts may become inconsistent if not performed atomically. Incorrect sign conventions for <code>quantity_change</code> may lead to misleading audit data.

3. UI/UX Design

Figure 3.1:




Figure 3.1.2:




Figure 3.2:

Welcome Back

Sign in to your account

Email *

Password *

SIGN IN

Figure 3.3:

Inventory Tracking System

HomeInventoryDashboardAdminAdmin

My Profile

Account Information

Name

Admin

Email

admin@hamfoodbank.ca

Role

Admin

Bank ID

1

Security

CHANGE PASSWORD

Figure 3.4:

Inventory Tracking System

HomeInventoryDashboardAdminAdmin

My Profile

Account Information

Name

Admin

Email

admin@hamfoodbank.ca

Role

Admin

Security

CHANGE PASSWORD

Change Password

Current Password *

New Password *

Confirm New Password *

CANCEL

UPDATE PASSWORD

Figure 3.5:




Figure 3.6:




Figure 3.7:




Figure 3.8:




Figure 3.9:




Figure 3.10:




Figure 3.11:

Inventory Tracking System

Home Inventory Dashboard **Admin** Admin

User Management

[+ CREATE USER](#)

NAME	EMAIL	ROLE	BANK ID	STATUS
Haley Johnson	haley@hamfoodbank.ca	Admin	1	Active
Tony Singh	tony@hamfoodbank.ca	Manager	1	Active
Marina Patel	marina@niagaraharvest.ca	Admin	2	Active
Liam Chen	liam@niagaraharvest.ca	Volunteer	2	Active
Isabella Brown	isabella@peelfamily.org	Admin	3	Active
David Moore	david@peelfamily.org	Volunteer	3	Active
Alex Rider	alex@hamfoodbank.ca	Driver	1	Active
Andy	Andy@hamfoodbank.ca	Admin	1	Active
Joe	Joe@hamfoodbank.ca	Manager	1	Active
Admin	admin@hamfoodbank.ca	Admin	1	Active

Figure 3.12:




Figure 3.13:

